

# assignment-4-21bit0376

September 22, 2023

NYSA SINGH

21BIT0376

```
[1]: #Importing essential libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: #Loading the dataset

df = pd.read_csv("/content/winequality-red.csv")
```

```
[4]: df.shape
```

```
[4]: (1599, 12)
```

```
[5]: df.head()
```

```
[5]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4             0.70         0.00           1.9       0.076
1           7.8             0.88         0.00           2.6       0.098
2           7.8             0.76         0.04           2.3       0.092
3          11.2             0.28         0.56           1.9       0.075
4           7.4             0.70         0.00           1.9       0.076
```

```
   free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0             11.0             34.0  0.9978  3.51       0.56
1             25.0             67.0  0.9968  3.20       0.68
2             15.0             54.0  0.9970  3.26       0.65
3             17.0             60.0  0.9980  3.16       0.58
4             11.0             34.0  0.9978  3.51       0.56
```

```
   alcohol  quality
0       9.4        5
1       9.8        5
2       9.8        5
```

|   |     |   |
|---|-----|---|
| 3 | 9.8 | 6 |
| 4 | 9.4 | 5 |

```
[18]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB

Checking for null values
```

```
[9]: df.isnull().any()
```

```
[9]: fixed acidity          False
volatile acidity          False
citric acid                False
residual sugar             False
chlorides                  False
free sulfur dioxide        False
total sulfur dioxide       False
density                    False
pH                         False
sulphates                  False
alcohol                    False
quality                    False
dtype: bool
```

We observe there are no null values, hence we move to data analysis step.

```
[10]: #Getting description of data
```

```
df.describe()
```

```
[10]:
```

|       | fixed acidity | volatile acidity | citric acid | residual sugar \ |
|-------|---------------|------------------|-------------|------------------|
| count | 1599.000000   | 1599.000000      | 1599.000000 | 1599.000000      |
| mean  | 8.319637      | 0.527821         | 0.270976    | 2.538806         |
| std   | 1.741096      | 0.179060         | 0.194801    | 1.409928         |
| min   | 4.600000      | 0.120000         | 0.000000    | 0.900000         |
| 25%   | 7.100000      | 0.390000         | 0.090000    | 1.900000         |
| 50%   | 7.900000      | 0.520000         | 0.260000    | 2.200000         |
| 75%   | 9.200000      | 0.640000         | 0.420000    | 2.600000         |
| max   | 15.900000     | 1.580000         | 1.000000    | 15.500000        |

|       | chlorides   | free sulfur dioxide | total sulfur dioxide | density \   |
|-------|-------------|---------------------|----------------------|-------------|
| count | 1599.000000 | 1599.000000         | 1599.000000          | 1599.000000 |
| mean  | 0.087467    | 15.874922           | 46.467792            | 0.996747    |
| std   | 0.047065    | 10.460157           | 32.895324            | 0.001887    |
| min   | 0.012000    | 1.000000            | 6.000000             | 0.990070    |
| 25%   | 0.070000    | 7.000000            | 22.000000            | 0.995600    |
| 50%   | 0.079000    | 14.000000           | 38.000000            | 0.996750    |
| 75%   | 0.090000    | 21.000000           | 62.000000            | 0.997835    |
| max   | 0.611000    | 72.000000           | 289.000000           | 1.003690    |


|       | pH          | sulphates   | alcohol     | quality     |
|-------|-------------|-------------|-------------|-------------|
| count | 1599.000000 | 1599.000000 | 1599.000000 | 1599.000000 |
| mean  | 3.311113    | 0.658149    | 10.422983   | 5.636023    |
| std   | 0.154386    | 0.169507    | 1.065668    | 0.807569    |
| min   | 2.740000    | 0.330000    | 8.400000    | 3.000000    |
| 25%   | 3.210000    | 0.550000    | 9.500000    | 5.000000    |
| 50%   | 3.310000    | 0.620000    | 10.200000   | 6.000000    |
| 75%   | 3.400000    | 0.730000    | 11.100000   | 6.000000    |
| max   | 4.010000    | 2.000000    | 14.900000   | 8.000000    |

```
[11]: #Checking values in each quality
df.quality.value_counts()
```

```
[11]: 5    681
      6    638
      7    199
      4     53
      8     18
      3     10
      Name: quality, dtype: int64
```

```
[12]: #Using catplot
sns.catplot(x='quality', data=df, kind='count')
```


```
[12]: <seaborn.axisgrid.FacetGrid at 0x7ea18a9b0880>
```



Checking for coloumn related to quality

```
[17]: sns.heatmap(df.corr(),annot=True, annot_kws={'size': 8}, fmt='.2f')
```

```
[17]: <Axes: >
```



```
[14]: df.corr()
```

```
[14]:
```

|                      | fixed acidity | volatile acidity | citric acid | \ |
|----------------------|---------------|------------------|-------------|---|
| fixed acidity        | 1.000000      | -0.256131        | 0.671703    |   |
| volatile acidity     | -0.256131     | 1.000000         | -0.552496   |   |
| citric acid          | 0.671703      | -0.552496        | 1.000000    |   |
| residual sugar       | 0.114777      | 0.001918         | 0.143577    |   |
| chlorides            | 0.093705      | 0.061298         | 0.203823    |   |
| free sulfur dioxide  | -0.153794     | -0.010504        | -0.060978   |   |
| total sulfur dioxide | -0.113181     | 0.076470         | 0.035533    |   |
| density              | 0.668047      | 0.022026         | 0.364947    |   |
| pH                   | -0.682978     | 0.234937         | -0.541904   |   |
| sulphates            | 0.183006      | -0.260987        | 0.312770    |   |
| alcohol              | -0.061668     | -0.202288        | 0.109903    |   |
| quality              | 0.124052      | -0.390558        | 0.226373    |   |

|               | residual sugar | chlorides | free sulfur dioxide | \ |
|---------------|----------------|-----------|---------------------|---|
| fixed acidity | 0.114777       | 0.093705  | -0.153794           |   |

|                      |           |           |           |
|----------------------|-----------|-----------|-----------|
| volatile acidity     | 0.001918  | 0.061298  | -0.010504 |
| citric acid          | 0.143577  | 0.203823  | -0.060978 |
| residual sugar       | 1.000000  | 0.055610  | 0.187049  |
| chlorides            | 0.055610  | 1.000000  | 0.005562  |
| free sulfur dioxide  | 0.187049  | 0.005562  | 1.000000  |
| total sulfur dioxide | 0.203028  | 0.047400  | 0.667666  |
| density              | 0.355283  | 0.200632  | -0.021946 |
| pH                   | -0.085652 | -0.265026 | 0.070377  |
| sulphates            | 0.005527  | 0.371260  | 0.051658  |
| alcohol              | 0.042075  | -0.221141 | -0.069408 |
| quality              | 0.013732  | -0.128907 | -0.050656 |

|                      | total sulfur dioxide | density   | pH        | sulphates | \ |
|----------------------|----------------------|-----------|-----------|-----------|---|
| fixed acidity        | -0.113181            | 0.668047  | -0.682978 | 0.183006  |   |
| volatile acidity     | 0.076470             | 0.022026  | 0.234937  | -0.260987 |   |
| citric acid          | 0.035533             | 0.364947  | -0.541904 | 0.312770  |   |
| residual sugar       | 0.203028             | 0.355283  | -0.085652 | 0.005527  |   |
| chlorides            | 0.047400             | 0.200632  | -0.265026 | 0.371260  |   |
| free sulfur dioxide  | 0.667666             | -0.021946 | 0.070377  | 0.051658  |   |
| total sulfur dioxide | 1.000000             | 0.071269  | -0.066495 | 0.042947  |   |
| density              | 0.071269             | 1.000000  | -0.341699 | 0.148506  |   |
| pH                   | -0.066495            | -0.341699 | 1.000000  | -0.196648 |   |
| sulphates            | 0.042947             | 0.148506  | -0.196648 | 1.000000  |   |
| alcohol              | -0.205654            | -0.496180 | 0.205633  | 0.093595  |   |
| quality              | -0.185100            | -0.174919 | -0.057731 | 0.251397  |   |

|                      | alcohol   | quality   |
|----------------------|-----------|-----------|
| fixed acidity        | -0.061668 | 0.124052  |
| volatile acidity     | -0.202288 | -0.390558 |
| citric acid          | 0.109903  | 0.226373  |
| residual sugar       | 0.042075  | 0.013732  |
| chlorides            | -0.221141 | -0.128907 |
| free sulfur dioxide  | -0.069408 | -0.050656 |
| total sulfur dioxide | -0.205654 | -0.185100 |
| density              | -0.496180 | -0.174919 |
| pH                   | 0.205633  | -0.057731 |
| sulphates            | 0.093595  | 0.251397  |
| alcohol              | 1.000000  | 0.476166  |
| quality              | 0.476166  | 1.000000  |

We observe that wine quality is highly influenced by alcohol and sulphates positively and presence of volatile acidity affects the quality of wine negatively.


Since we are performing classification and we know that Classification Algorithm like decision tree and Random forest are not affected by the UnScaled data. Hence we are Avoiding scaling of data.

Classification Algorithm are also insensitive to outliers, but as a part of data preprocessing we will

try to detect and remove outliers.

```
[19]: #Checking for fixed Acidity  
sns.boxplot(df['fixed acidity'])
```

```
[19]: <Axes: >
```




We Observe there are outliers in the fixed Acidity columns. Hence, we will remove them.

```
[27]: #Using IQR method  
  
f1= df['fixed acidity'].quantile(0.25) #First Quartile  
f3= df['fixed acidity'].quantile(0.75) #Third Quartile  
IQR=f3-f1 #Inter Quartile range  
  
Upper_limit = f3+(1.5)*IQR  
Lower_limit = f1-(1.5)*IQR  
  
df=df[(df['fixed acidity']<Upper_limit) & (df['fixed acidity']>Lower_limit)]
```

```
[26]: sns.boxplot(df['fixed acidity'])
```

[26]: <Axes: >




checking for outliers in volatile acidity column

```
[28]: sns.boxplot(df['volatile acidity'])
```

[28]: <Axes: >





We see outliers are present in Volatile Acidity column. Hence we will remove the outliers

```
[29]: #Using IQR method


f1= df['volatile acidity'].quantile(0.25) #First Quartile
f3= df['volatile acidity'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['volatile acidity']<Upper_limit) & (df['volatile_
↪acidity']>Lower_limit)]
```

```
[30]: sns.boxplot(df['volatile acidity'])
```


```
[30]: <Axes: >
```



Checking for outliers in citric Acid column

```
[31]: sns.boxplot(df['citric acid'])
```

```
[31]: <Axes: >
```



We observe there is a outlier in citric Acid column. Hence we will remove it.

```
[32]: #Using IQR method


f1= df['citric acid'].quantile(0.25) #First Quartile
f3= df['citric acid'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['citric acid']<Upper_limit) & (df['citric acid']>Lower_limit)]
```

```
[33]: sns.boxplot(df['citric acid'])
```

```
[33]: <Axes: >
```




We see that Residual sugar does not affect Quality much. Hence we skip the Residual Sugar Column.

Checking for Outliers in chlorides

```
[34]: sns.boxplot(df['chlorides'])
```

```
[34]: <Axes: >
```



We observe there are a large number of outliers in chlorides columns. Hence removing them.

[39]: *#Using IQR method*


```
f1= df['chlorides'].quantile(0.25) #First Quartile
f3= df['chlorides'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range
```

```
Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR
```

```
df=df[(df['chlorides']<Upper_limit) & (df['chlorides']>Lower_limit)]
```

[40]: sns.boxplot(df['chlorides'])

[40]: <Axes: >




We see that Free Sulphur Dioxide does not affect Quality much. Hence we skip the Free sulphur dioxide Column.

Checking for Outliers in Total sulphur dioxide

```
[41]: sns.boxplot(df['total sulfur dioxide'])
```

```
[41]: <Axes: >
```



We observe there are a large number of outliers in Total sulphur Dioxide columns. Hence removing them.

[46]: *#Using IQR method*


```
f1= df['total sulfur dioxide'].quantile(0.25) #First Quartile
f3= df['total sulfur dioxide'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['total sulfur dioxide']<Upper_limit) & (df['total sulfur_
    dioxide']>Lower_limit)]
```

[47]: `sns.boxplot(df['total sulfur dioxide'])`

[47]: <Axes: >




Checking for Outliers in Density

```
[48]: sns.boxplot(df.density)
```

```
[48]: <Axes: >
```





We observe there are outliers in Density columns. Hence removing them.

```
[51]: #Using IQR method


f1= df['density'].quantile(0.25) #First Quartile
f3= df['density'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['density']<Upper_limit) & (df['density']>Lower_limit)]
```

```
[52]: sns.boxplot(df.density)
```

```
[52]: <Axes: >
```




We see that PH does not affect Quality much. Hence we skip the PH Column.

Checking for outliers in Sulphates column

```
[53]: sns.boxplot(df.sulphates)
```

```
[53]: <Axes: >
```



We observe there are a large number of outliers in Sulphates columns. Hence removing them.

[56]: *#Using IQR method*


```
f1= df['sulphates'].quantile(0.25) #First Quartile
f3= df['sulphates'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['sulphates']<Upper_limit) & (df['sulphates']>Lower_limit)]
```

[57]: sns.boxplot(df.sulphates)


[57]: <Axes: >



Checking for outliers in Alcohol Column

```
[58]: sns.boxplot(df.alcohol)
```

```
[58]: <Axes: >
```



There are outliers in Alcohol. Hence removing them.

```
[59]: #Using IQR method


f1= df['alcohol'].quantile(0.25) #First Quartile
f3= df['alcohol'].quantile(0.75) #Third Quartile
IQR=f3-f1 #Inter Quartile range

Upper_limit = f3+(1.5)*IQR
Lower_limit = f1-(1.5)*IQR

df=df[(df['alcohol']<Upper_limit) & (df['alcohol']>Lower_limit)]
```

```
[61]: sns.boxplot(df.alcohol)
```

```
[61]: <Axes: >
```



We Observe that outliers in all the Columns are removed.

Seperating data into dependent and independent columns

```
[63]: X= df.drop('quality',axis=1);
```

```
[64]: X
```

```
[64]:
```

|      | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides \ |
|------|---------------|------------------|-------------|----------------|-------------|
| 0    | 7.4           | 0.700            | 0.00        | 1.9            | 0.076       |
| 1    | 7.8           | 0.880            | 0.00        | 2.6            | 0.098       |
| 2    | 7.8           | 0.760            | 0.04        | 2.3            | 0.092       |
| 3    | 11.2          | 0.280            | 0.56        | 1.9            | 0.075       |
| 4    | 7.4           | 0.700            | 0.00        | 1.9            | 0.076       |
| ...  | ...           | ...              | ...         | ...            | ...         |
| 1594 | 6.2           | 0.600            | 0.08        | 2.0            | 0.090       |
| 1595 | 5.9           | 0.550            | 0.10        | 2.2            | 0.062       |
| 1596 | 6.3           | 0.510            | 0.13        | 2.3            | 0.076       |
| 1597 | 5.9           | 0.645            | 0.12        | 2.0            | 0.075       |
| 1598 | 6.0           | 0.310            | 0.47        | 3.6            | 0.067       |

|   | free sulfur dioxide | total sulfur dioxide | density | pH   | sulphates \ |
|---|---------------------|----------------------|---------|------|-------------|
| 0 | 11.0                | 34.0                 | 0.99780 | 3.51 | 0.56        |

|      |      |      |         |      |      |
|------|------|------|---------|------|------|
| 1    | 25.0 | 67.0 | 0.99680 | 3.20 | 0.68 |
| 2    | 15.0 | 54.0 | 0.99700 | 3.26 | 0.65 |
| 3    | 17.0 | 60.0 | 0.99800 | 3.16 | 0.58 |
| 4    | 11.0 | 34.0 | 0.99780 | 3.51 | 0.56 |
| ...  | ...  | ...  | ...     | ...  | ...  |
| 1594 | 32.0 | 44.0 | 0.99490 | 3.45 | 0.58 |
| 1595 | 39.0 | 51.0 | 0.99512 | 3.52 | 0.76 |
| 1596 | 29.0 | 40.0 | 0.99574 | 3.42 | 0.75 |
| 1597 | 32.0 | 44.0 | 0.99547 | 3.57 | 0.71 |
| 1598 | 18.0 | 42.0 | 0.99549 | 3.39 | 0.66 |

|      | alcohol |
|------|---------|
| 0    | 9.4     |
| 1    | 9.8     |
| 2    | 9.8     |
| 3    | 9.8     |
| 4    | 9.4     |
| ...  | ...     |
| 1594 | 10.5    |
| 1595 | 11.2    |
| 1596 | 11.0    |
| 1597 | 10.2    |
| 1598 | 11.0    |

[1184 rows x 11 columns]

```
[65]: #Getting Dependent variable
      Y=df.quality
```

```
[66]: Y
```

```
[66]: 0      5
      1      5
      2      5
      3      6
      4      5
      ..
      1594    5
      1595    6
      1596    6
      1597    5
      1598    6
      Name: quality, Length: 1184, dtype: int64
```

Performing Binarization

We are considering quality above 6 means wine is good denoted as 1 else its bad denoted as 0

```
[67]: Y=Y.apply(lambda y_value :1 if y_value>=7 else 0)
```

```
[68]: Y
```

```
[68]: 0      0
      1      0
      2      0
      3      0
      4      0
      ..
     1594    0
     1595    0
     1596    0
     1597    0
     1598    0
      Name: quality, Length: 1184, dtype: int64
```

### Performing train Test split

```
[62]: from sklearn.model_selection import train_test_split
```

```
[92]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.
      ↪3,random_state=4)
```

```
[93]: print(df.shape, X_train.shape, X_test.shape, Y_train.shape, Y_test.shape)

(1184, 12) (828, 11) (356, 11) (828,) (356,)
```

### Using Random Forest Classifier

```
[81]: from sklearn.ensemble import RandomForestClassifier

      model = RandomForestClassifier()
```

```
[94]: #Fitting Data

      model.fit(X_train,Y_train)
```

```
[94]: RandomForestClassifier()
```

```
[95]: y_predict = model.predict(X_test)

      y_predict_train = model.predict(X_train)
```

### Model Evaluation

```
[98]: from sklearn.metrics import accuracy_score
```

```
[96]: print("Accuracy: ", accuracy_score(Y_test,y_predict))
```



Accuracy: 0.9297752808988764

```
[97]: print("Accuracy: ", accuracy_score(Y_train,y_predict_train))
```

Accuracy: 1.0

### Testing with Random Values

```
[104]: data = (8.3,      0.84,      0.07,      1.9,      0.1,      18.  
             ↪0,      43,      0.884,      4.56,      0.87,      8.9)  
model.predict((np.asarray(data)).reshape(1,-1))
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature names

```
warnings.warn(
```

```
[104]: array([0])
```

We got label as 0, Means the wine is of Bad Quality